

R Notebook - Billboard 100 Data Cleaning & Visualization

Maya Reese Farmer

October 16, 2021

Introduction

I will be demonstrating some data cleaning and visualization techniques using the dplyr and tidyr packages in R. I will also be using the Billboard Hot 100 dataset from kaggle.com

This dataset contains the following variables:

- rank
- song
- artist
- last-week (rank in previous week)
- peak-week (top rank achieved by song)
- weeks-on-board
- date

```
## importing dataset into r
charts <- read.csv("~/Downloads/charts.csv")
attach(charts)
head(charts)
```

```
##   rank      song      artist last.week
## 1    1    Butter      BTS         1
## 2    2    Good 4 U  Olivia Rodrigo   2
## 3    3 Levitating Dua Lipa Featuring DaBaby   4
## 4    4 Kiss Me More Doja Cat Featuring SZA   3
## 5    5 Montero (Call Me By Your Name) Lil Nas X   8
## 6    6 Bad Habits Ed Sheeran   5
##   peak.rank weeks.on.board      date
## 1         1             7 2021-07-17
## 2         1             8 2021-07-17
## 3         2            40 2021-07-17
## 4         3            13 2021-07-17
## 5         1            15 2021-07-17
## 6         5             2 2021-07-17
```

Data Cleaning

```
##  
library(dplyr)
```

task 1: convert date column to date format

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(lubridate)
```

```
##  
## Attaching package: 'lubridate'  
  
## The following objects are masked from 'package:base':  
##  
##   date, intersect, setdiff, union
```

```
library(stringr)
```

```
charts %>% select(date, song, artist) %>% mutate(date = ymd(date)) %>% head(10)
```

```
##           date                song  
## 1  2021-07-17                Butter  
## 2  2021-07-17                Good 4 U  
## 3  2021-07-17            Levitating  
## 4  2021-07-17            Kiss Me More  
## 5  2021-07-17 Montero (Call Me By Your Name)  
## 6  2021-07-17            Bad Habits  
## 7  2021-07-17    Leave The Door Open  
## 8  2021-07-17            Peaches  
## 9  2021-07-17    Save Your Tears  
## 10 2021-07-17            Deja Vu  
##           artist  
## 1                BTS  
## 2            Olivia Rodrigo  
## 3    Dua Lipa Featuring DaBaby  
## 4    Doja Cat Featuring SZA  
## 5                Lil Nas X  
## 6            Ed Sheeran  
## 7    Silk Sonic (Bruno Mars & Anderson .Paak)  
## 8  Justin Bieber Featuring Daniel Caesar & Giveon  
## 9            The Weeknd & Ariana Grande  
## 10            Olivia Rodrigo
```

```
charts %>% mutate(date = ymd(date)) %>%
  distinct(date) %>%
  mutate(month = floor_date(date, "month"), year = floor_date(date, "year")) %>% head(10)
```

task 2: create 2 new columns containing month and year

```
##           date      month      year
## 1  2021-07-17 2021-07-01 2021-01-01
## 2  2021-07-10 2021-07-01 2021-01-01
## 3  2021-07-03 2021-07-01 2021-01-01
## 4  2021-06-26 2021-06-01 2021-01-01
## 5  2021-06-19 2021-06-01 2021-01-01
## 6  2021-06-12 2021-06-01 2021-01-01
## 7  2021-06-05 2021-06-01 2021-01-01
## 8  2021-05-29 2021-05-01 2021-01-01
## 9  2021-05-22 2021-05-01 2021-01-01
## 10 2021-05-15 2021-05-01 2021-01-01
```

task 3: separate the artist column Here, I want to break up the artist column so that each artist is designated as a primary artist (the first artist) or a featured artist (the 2nd, 3rd, or 4th artist) in different columns. To do that, I'm going to use the separate function in the tidyr package. With this I'm able to designate what characters or values I want to use as separators between columns and I can also name columns. Then, I'll create a new data frame called clean.charts that contains the original data along with the updated date columns and artist columns.

```
library(tidyr)

clean.charts = charts %>%
  mutate(date = ymd(date)) %>%
  mutate(month = floor_date(date, "month"), year = floor_date(date,
    "year")) %>%
  separate(artist, c("primary.artist", "featured.artist", "featured.artist2",
    "featured.artist3"), "\\sFeaturing\\s|\\s&\\s|\\s|\\sX\\s|\\sx\\s|\\sWith\\s",
    remove = FALSE)
```

```
## Warning: Expected 4 pieces. Additional pieces discarded in 665 rows [888, 943,
## 3392, 3490, 3591, 3679, 3773, 3795, 3875, 3897, 3992, 4076, 4100, 4176, 4198,
## 4275, 4299, 4364, 4398, 4442, ...].
```

```
## Warning: Expected 4 pieces. Missing pieces filled with 'NA' in 326334 rows [1,
## 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, ...].
```

```
clean.charts %>%
  select(song, primary.artist, featured.artist, featured.artist2,
    featured.artist3) %>%
  head(10)
```

```
##           song      primary.artist featured.artist
## 1      Butter                BTS                <NA>
```

```
## 2          Good 4 U          Olivia Rodrigo          <NA>
## 3          Levitating          Dua Lipa          DaBaby
## 4          Kiss Me More          Doja Cat          SZA
## 5  Montero (Call Me By Your Name)          Lil Nas X          <NA>
## 6          Bad Habits          Ed Sheeran          <NA>
## 7          Leave The Door Open Silk Sonic (Bruno Mars Anderson .Paak)
## 8          Peaches          Justin Bieber          Daniel Caesar
## 9          Save Your Tears          The Weeknd          Ariana Grande
## 10         Deja Vu          Olivia Rodrigo          <NA>
##   featured.artist2 featured.artist3
## 1          <NA>          <NA>
## 2          <NA>          <NA>
## 3          <NA>          <NA>
## 4          <NA>          <NA>
## 5          <NA>          <NA>
## 6          <NA>          <NA>
## 7          <NA>          <NA>
## 8          Giveon          <NA>
## 9          <NA>          <NA>
## 10         <NA>          <NA>
```

Overall, this is a pretty good way to separate these columns, but it isn't perfect. The original artist column is very messy and artists aren't separated the same way. For example, one of the separators I had to use was a comma. For artists that have a comma in their name, like Tyler, The Creator, our code will break up his name into Tyler in the primary artist column and The Creator in the featured artist column. This is something to remember if you're trying to isolate a single artist for further analysis later. But for the most part, the separate function did a pretty good job of breaking up the data into appropriate columns.

task 4: pivot table Now that the data are fairly clean, I want to collapse the primary artist and featured artist columns (total 4 columns) into just 2 columns using the `pivot_longer` function in `tidyr`. This format will be more conducive for analyses in the future.

```
pivot.artist = clean.charts %>%
  pivot_longer(primary.artist:featured.artist3, names_to = "artist.type",
               values_to = "artist.name")
pivot.artist %>%
  select(song, artist.type, artist.name) %>%
  head(15)
```

```
## # A tibble: 15 x 3
##   song          artist.type  artist.name
##   <chr>         <chr>         <chr>
## 1 Butter       primary.artist  BTS
## 2 Butter       featured.artist <NA>
## 3 Butter       featured.artist2 <NA>
## 4 Butter       featured.artist3 <NA>
## 5 Good 4 U     primary.artist  Olivia Rodrigo
## 6 Good 4 U     featured.artist <NA>
## 7 Good 4 U     featured.artist2 <NA>
## 8 Good 4 U     featured.artist3 <NA>
## 9 Levitating   primary.artist  Dua Lipa
## 10 Levitating  featured.artist  DaBaby
## 11 Levitating  featured.artist2 <NA>
```

```
## 12 Levitating    featured.artist3 <NA>
## 13 Kiss Me More  primary.artist    Doja Cat
## 14 Kiss Me More  featured.artist    SZA
## 15 Kiss Me More  featured.artist2 <NA>
```

Now that we have a new table where the artist name and type have been pivoted/collapsed into 2 columns, we can easily query artists and determine the number of songs where they were features or the primary artist.

```
## example using SZA
pivot.artist %>% distinct(song, artist.type,artist.name) %>%
  filter(artist.name == "SZA")
```

```
## # A tibble: 14 x 3
##   song                artist.type    artist.name
##   <chr>              <chr>      <chr>
## 1 Kiss Me More      featured.artist SZA
## 2 Good Days        primary.artist SZA
## 3 Hit Different     primary.artist SZA
## 4 The Other Side    primary.artist SZA
## 5 Staring At The Sun featured.artist SZA
## 6 Just Us           featured.artist SZA
## 7 Power Is Power    primary.artist SZA
## 8 All The Stars     featured.artist SZA
## 9 I Do              featured.artist SZA
## 10 Broken Clocks    primary.artist SZA
## 11 What Lovers Do   featured.artist SZA
## 12 The Weekend      primary.artist SZA
## 13 Love Galore      primary.artist SZA
## 14 Homemade Dynamite featured.artist3 SZA
```

task 5: join genre Finally, I want to join another data frame, the `mtv10000` data, to my `clean.charts` data frame. The `mtv10000` dataset will provide the genre of music for each primary artist. Because I want to retain all the rows in my ‘clean.charts’ data frame, I’m going to perform a left join. A left join keeps all data in the left table (`clean.charts`) and only the rows that match the condition in the second table will be added (NAs will be added in rows where there is no match in the second table). Essentially, if an artist in the `clean.charts` table is also in the `mtv10000` table, then a genre will be returned. If an artist in the `clean.charts` table is not in the `mtv10000` table, then a NA will be returned.

```
mtv10000 <- read.csv("~/Downloads/mtv10000.csv")
attach(mtv10000)

## it's VERY important to trim the white spaces from the
## data before performing any joins
clean.charts$primary.artist = trimws(clean.charts$primary.artist)
mtv10000$name = trimws(mtv10000$name)
attach(clean.charts)
```

```
## The following objects are masked from charts:
##
##   artist, date, last.week, peak.rank, rank, song, weeks.on.board
```

```
attach(mtv10000)
```

```
## The following objects are masked from mtv10000 (pos = 4):  
##  
##    facebook, genre, mtv, name, twitter, website
```

```
clean.charts %>%  
  left_join(mtv10000, by = c(primary.artist = "name")) %>%  
  select(primary.artist, genre) %>%  
  head(10)
```

	primary.artist	genre
## 1	BTS	<NA>
## 2	Olivia Rodrigo	<NA>
## 3	Dua Lipa	<NA>
## 4	Doja Cat	<NA>
## 5	Lil Nas X	<NA>
## 6	Ed Sheeran	Singer/Songwriter
## 7	Silk Sonic (Bruno Mars	<NA>
## 8	Justin Bieber	Pop
## 9	The Weeknd	R&B/Soul
## 10	Olivia Rodrigo	<NA>

So, the left join was successful, but as you can see there are a lot of missing artists in the mtv10000 dataset. This means that there are a lot of missing data points in the genre column of our joined data. If this were for a real project, the mtv10000 data would not be the best joining genre because the lack of data might skew our analyses. But it's fine for this informal project.

Now, we can use our cleaned and joined data to create a few visualizations!

Data Visualization

First, let's look at how the popularity of the top 5 genres have changed over time.

```
## create new dataframe from join  
genre.join = clean.charts %>% left_join(mtv10000, by = c("primary.artist" = "name"))  
  
## find the top 5 genres in df  
top5.genre = genre.join %>% select(date, primary.artist, genre) %>%  
  filter(!is.na(genre)) %>%  
  distinct() %>%  
  count(genre) %>%  
  top_n(5) %>%  
  pull(genre)
```

```
## Selecting by n
```

```
top5.genre
```

```
## [1] "Country"      "Hip-Hop/Rap" "Pop"          "R&B/Soul"     "Rock"

## create new dataframe only including yearly counts of top
## 5 genres
genre.year = genre.join %>%
  select(year, song, genre) %>%
  filter(genre %in% top5.genre) %>%
  count(year, genre)

## load ggplot2 and cowplot for graphs
library(ggplot2)
library(cowplot)

##
## Attaching package: 'cowplot'


## The following object is masked from 'package:lubridate':
##
## stamp

library(wesanderson)

cavalcanti = wes_palette("Cavalcanti1")

ggplot(genre.year, aes(year, n, color = genre)) + geom_line(size = 1) +
  theme_cowplot(12) + scale_color_manual(values = cavalcanti,
    name = "Genre") + labs(x = "Year", y = "Number of Songs on Billboard 100",
    title = "Popularity of Top 5 Genres over Time") + theme(plot.title = element_text(size = 16,
    face = "bold", hjust = 0.5), axis.text.x = element_text(size = 12),
    axis.title = element_text(face = "bold"), axis.text.y = element_text(size = 12),
    legend.title = element_text(face = "bold"))
```

Popularity of Top 5 Genres over Time



It's very important to note that I cannot draw any concrete conclusions based on this figure because of the missing genre data from our join. It's clear that the mtv10000 data was missing new artists from the past 10 years because there is an overall decline in all genres starting around 2010 to the present (based on the figure). However, it's still pretty interesting to see how the popularity of the top 5 genres have changed over time. For example, the popularity of country music increased significantly in the early 2000s. Pop music had the highest popularity overall around 2010. Also, Hip-Hop and Rap was a new genre on the Billboard 100 around 1980, but it is still in the top 5 genres in the dataset (very popular).

Next, let's look at the top 5 songs that were repeated ranked in the top 10 in 2019.

```
top5.rank = clean.charts %>%
  select(rank, song, year) %>% filter(year == "2019-01-01", rank < 11) %>%
  count(song) %>% arrange(desc(n)) %>% top_n(5)
```


Selecting by n

```
top5.rank
```

```
##           song  n
## 1           Bad Guy 30
## 2 Sunflower (Spider-Man: Into The Spider-Verse) 30
## 3           Old Town Road 26
## 4                Wow. 24
## 5           No Guidance 23
## 6           Senorita 23
```



```
ggplot(top5.rank, aes(x = reorder(song, n), y = n, fill = reorder(song,
n), color = reorder(song, n))) + geom_bar(stat = "identity",
size = 1) + coord_flip() + theme_cowplot(12) + labs(x = "",
y = "", title = "Most Popular Songs in 2019") + scale_fill_manual(values = alpha(c("#1A8693",
"#4D7A5C", "#D4B95E", "#D28F33", "#B34233", "#751A33"), 0.7)) +
scale_color_manual(values = c("#1A8693", "#4D7A5C", "#D4B95E",
"#D28F33", "#B34233", "#751A33")) + theme(plot.title = element_text(size = 16,
face = "bold", hjust = 0.5), legend.position = "none", axis.text = element_text(face = "bold",
size = 11)) + scale_x_discrete(labels = c("No Guidance",
"Senorita", "Wow.", "Old Town Road", "Bad Guy", "Sunflower (Spider-Man:
Into The Spider-Verse)"))
```



This shows the songs that repeatedly ranked in the top 10 (most popular) on the Billboard 100 in 2019. For example, Old Town Road ranked in the top 10 for 26 weeks in 2019.

Finally, I'm a Beyonce fan so I want to explore her top 10 most popular songs on Billboard 100.

```
beyonce10 = clean.charts %>%
  select(primary.artist, song, month) %>%
  filter(primary.artist == "Beyonce") %>%
  distinct() %>%
  count(song) %>%
  arrange(desc(n)) %>%
  head(10)
beyonce10
```

```
##              song n
## 1              Halo 8
## 2          Baby Boy 7
```

```
## 3          Check On It 7
## 4          Crazy In Love 7
## 5          Irreplaceable 7
## 6          Love On Top 7
## 7 Single Ladies (Put A Ring On It) 7
## 8          Sweet Dreams 7
## 9          Upgrade U 7
## 10         ***Flawless 6
```

```
ggplot(beyonce10, aes(reorder(song, -n), n)) + geom_segment(aes(x = reorder(song,
-n), xend = reorder(song, n), y = 0, yend = n), color = "grey") +
  geom_point(size = 5, color = "steelblue", fill = alpha("steelblue",
0.3), alpha = 0.7, shape = 21, stroke = 2) + theme_cowplot(12) +
  labs(x = "", y = "Number of Times on Billboard 100", title = "Most Popular Beyonce Songs") +
  scale_x_discrete(labels = c("Halo", "Baby\n Boy", "Check\n On It",
"Crazy\n In Love", "Irreplaceable", "Love\n On Top", "Single\n Ladies",
"Sweet\n Dreams", "Upgrade\n U", "Flawless")) + theme(panel.grid.major = element_blank(),
panel.grid.minor = element_blank(), panel.background = element_blank(),
plot.title = element_text(size = 16, face = "bold", hjust = 0.5),
axis.ticks = element_blank(), axis.text.x = element_text(size = 10),
axis.text.y = element_text(size = 11), axis.line = element_blank()) +
  ylim(0, 9)
```

Most Popular Beyonce Songs

